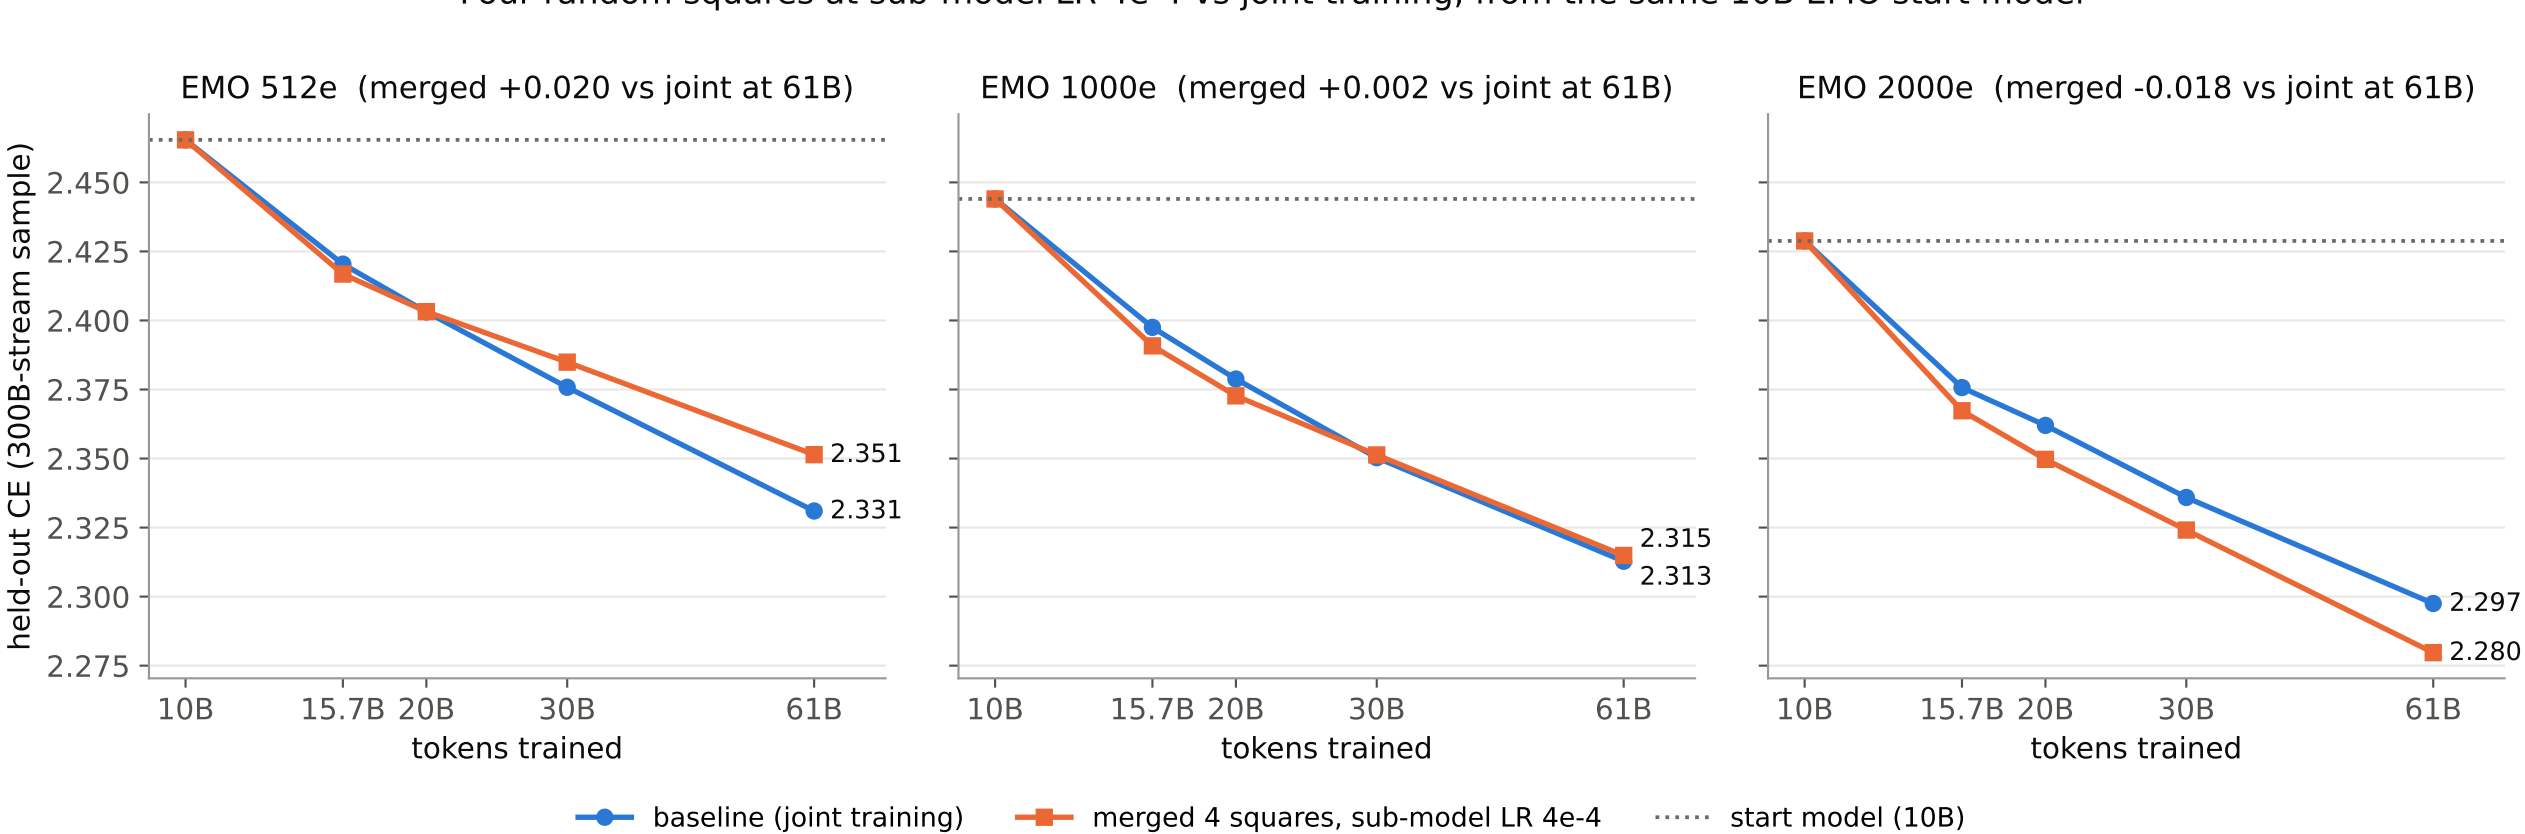
Four random squares at sub-model LR 4e-4 vs joint training, from the same 10B EMO start model
EMO 512e (merged +0.020 vs joint at 61B)
EMO 1000e (merged +0.002 vs joint at 61B)
EMO 2000e (merged -0.018 vs joint at 61B)
held-out CE (300B-stream sample)
2.450
2.425
2.400
2.375
2.350
2.325
2.300
2.275
2.351
2.331
2.315
2.313
2.297
2.280
10B
15.7B
20B
30B
61B
tokens trained
baseline (joint training)
merged 4 squares, sub-model LR 4e-4
start model (10B)